

Supervised Machine Learning HW 2

Group 3: Maja Nordfeldt, Jakob Rauch, Timo Schenk, Konstantin Sommer

16-1-2020

Introduction

Which demographic characteristics predict turnover of grocery stores? A high turnover is an important factor when supermarkets choose the location of a new brand. Especially in urban areas, data on a large number of demographic factors is available. From an intuitive perspective, most of the variables might play a role in determining sales. However, when predicting the turnover from data, it is often better to rely only on a small subset of the variables. Using data from supermarkets in the Chicago area, we estimate an elastic net and find 8 relevant predictors out of 44 observed demographic factors. We conclude that a prediction of turnovers in a certain location should rely on the share of small households, average age, share of high income households and unemployment measures.

Data

The data is store-level scanner data collected at 77 supermarkets in the Chicago area during 1996, in a collaboration between Chicago Booth School of Business and Dominick's Finer Foods¹. Store data is matched with demographic census data of the metropolitan Chicago area, originally obtained from the US Government in 1990.

The dependent variable *grocery_sum* is the total turnover in one year of groceries, measured in \$. The predictor variables present the demographics of the respective stores nearby area - *age9* and *age60* represent the percent of population under age 9 and 60 respectively. *ethnic* presents percent of blacks and hispanics, *nwhite* indicates percent of population that is non-white, *educ* percent of college graduates, *ncar* percent with no vehicles, *income* the log of the median income, *poverty* indicates percent of population with income below \$15,000, *incsigma* the approximated standard deviation of the income distribution, *hsizeavg* the average household size, *hsize_x* and *hh_xplus* variables denote percent of households containing *x* persons, *hhsingle* measures percentage of detached houses, *hhlarge* the percentage with 5 or more people, *workwom* percentage of working women with full time jobs, *sinhouse* the percentage of one-person households, *density* the trading area in square miles per capita, *hval_x* the percent of households with value over *x* thousands of dollars, *hvalmean* the approximated mean household value, *single* the percentage of singles, *retired** the percentage of retired and *unemp* the percentage of unemployed. *wrk* variables denote percent of women working, with a *ch* ending denoting having children, *nch* no children and *nwrk* not working, and the number indicating age of children (5 means children younger than 5 years old and 17 means children are between 6 - 17 years old). *telephn* and *mortgage* indicates percent of households having a telephone and mortgage respectively. Further, *shop_x* denotes percent of type *x* of shopper - including constrained, hurried, avid, unfettered, birds and stranges. Finally, *shopindx* and *shpinx** present ability to shop, meaning having a car and being a single family house.

Note that from the original dataset we transform all variables to z-scores, removing their unit of measure. After this transformation, each variable's value can be interpreted as standard deviations from its mean. This was undertaken to avoid the original magnitude of the variables influencing their estimates.

¹The dataset can be retrieved from <https://www.chicagobooth.edu/research/kilts/datasets/dominicks>

Methods

Since we are dealing with a dataset with about 44 predictors and only 77 observations we use a methodology that induces both shrinkage of parameter estimates and variable selection. This is because a situation in which the ratio of observations to predictor variables is low, there is a danger of what is called “overfitting”. In such a situation we might be able to find unbiased estimates of a coefficient for each of the variables using the standard multiple regression framework but since there are only few observations available per parameter, the model will do very poorly in out of sample prediction. We therefore use a combination of Ridge and Lasso regressions called elastic net. Both of these methods, as well as their combination, are aimed at driving some of the coefficients towards zero such that they are not used in prediction anymore. In order to do so we extend the loss function, which in OLS is only the residuals sum of square, by a penalty term, that increases with the size of the coefficients. So when minimizing the loss function one can also think about the penalty term as a constraint on the size of the coefficients. Since both of these methods have advantages and disadvantages, which we will describe in the following, we use a combination of the two, the elastic net method.

The loss function for ridge regression is $(y - X\beta)^T(y - X\beta) + \lambda\beta^T\beta$. The first term is the sum of squared residuals, so the loss function of the multiple regression setting. As one can see we are here not only interested in minimizing this but in addition we also add the penalty term $\lambda\beta^T\beta$, where λ is the penalty strength. One can see that the larger the chosen value of λ the lower will be the optimal coefficients. The penalty term increases here quadratically in β , which is the key difference to the LASSO method in which the penalty term is $\lambda\|\beta\|_1$, where $\|\beta\|_1$ refers to the L1 norm so the sum of the absolute values of the vector. So the penalty term in LASSO increases with the absolute value of the coefficients and one can again state that the larger the chosen penalty term the smaller will be the resulting coefficient values.

The main difference between the two penalty terms is that while Ridge drives the coefficient values close to zero but mostly not exactly to zero, LASSO will assign an exact zero to many of the coefficients. Why this is the case becomes a bit clearer when looking at how the penalty term translate into a constraint on minimizing the RSS. While for Ridge this is $\beta^T\beta < \gamma$, for LASSO this is $\|\beta\|_1 < \gamma$. Both of these constraints are equivalent to minimizing the respective loss function and both span a space around the origin in which the optimal coefficient values must be. While in Ridge this is a hyperbole, in LASSO this is a space with straight edges. One can see that when analysing the two dimensional case that it is mostly the case that for LASSO one of the coefficients will be zero and the other one will be relatively large, for Ridge, both coefficients are driven towards zero but not to exactly zero.

Driving some of the coefficients exactly to zero has the advantage that is basically translates into a variable selection exercise, since these coefficients will then not be used in prediction anymore. On the other hand LASSO might also drive coefficients towards zero that are actually not zero and that are then excluded eventhough they might actually explain some of the variation of the endogenous variable. So to use the advantages of both methods we will combine the two in an elastic net approach in which the penalty term is set up as a combination of the two previous ones: $\lambda(\alpha\|\beta\|_1 + (1 - \alpha)\beta^T\beta)$, where λ still determines the penalty strength and α determines the importance of the LASSO penalty relative to the Ridge penalty.

Any deviation from the OLS loss function will result in biased estimates for the coefficients, however, a penalty term leads to the resulting coefficients exhibiting smaller variances. Additionally, OLS does not work in the presence of exact or strong multicollinearity since the X matrix of predictors will become singular or close to singular, making the inversion of the $X'X$ matrix impossible or very inexact (due to many rounding errors). More importantly, a sufficiently high penalty prevents overfitting and, in the presence of many predictors and few observations, will outperform a multiple regression estimation in predicting out of sample values. One therefore faces what is known as a bias variance trade of, determined by the penalty strength λ .

Using the elastic net method, we need to decide for values of the two hyperparameters: the penalty λ and α , the weight of LASSO relative to Ridge regression. There is an optimal combination of values for these hyperparameters, in the sense of delivering the best out-of-sample performance. To find this optimal combination, we use a method called k-fold cross-validation. In this procedure, we first arrange our dataset into random order (as the order of observations might be correlated with some features of the data, e.g. when

neighborhoods are sampled in a spatial order. We then split the data into k bins of equal size. We use the last $(k-1)$ bins as a training sample (i.e. we fit our model to this data and end up with parameter estimates) and use the first bin as test sample, where we calculate the out-of-sample mean squared error, a measure of the distance between predicted values given our parameter estimates and the observed values. We repeat this k times, using in the second/third/... iteration the second/third/... bin for testing and the remaining bins for training. We finally take the square root of the average of these k mean squared errors, which will be our measure of model performance given the hyperparameters. To find the optimal values, we do a grid search - calculating the model performance for each combination of α and λ in the grid. The best values will lead to the lowest average mean squared error. These values lead to a model with good predictive performance.

Results

Table 1: Non-zero point estimates for $(\lambda, \alpha) = (0.04, 0.6)$

AGE9	HSIZE2	HSIZE34	SINHOUSE	HVAL150	HVAL200	UNEMP	SHPHURR
0.3759	0.02787	0.1159	0.1104	0.1978	0.01314	0.254	0.1544

We applied the method to the Chicago data to select the relevant predictors out of 44 observed demographic factors. Table 1 shows the estimates of the 8 coefficients that the elastic net has not shrunked to zero. The optimal shrinkage parameter was obtained by K-fold crossvalidation, splitting the data into $k = 7$ parts to get equally large bins for 77 observations. Our manually programmed results are close to those from the R-package `glmnet()`. Yet the estimates are not exactly the same, as both `glmnet()` and our manual function have non-deterministic results because data is randomly rearranged. Best values are $\alpha = 0.6$ and $\lambda = 0.04$ for a grid search with α ranging between 0 and 1, with stepsize 0.1, and λ ranging between 0.01 and 1 with stepsize 0.01.

Conclusion

In this study we aimed to identify what demographics that are linked to higher store turnovers. XXX[ALTER TO RQ formulation] sales. Using an elastic-net strategy, we identified a few key predictors.

Firstly, it appears that the household size is relevant - in particular single and three to four-person households have explanatory power. Even more important appears the share of high-income household values to be, in particular a household value over 150 000 dollar. The return on targeting a neighbourhood with a higher share of value over 200 000 dollar is however not as large. Somewhat contradictory to the income predictors is the share of unemployed, being the single highest predictor. Finally, hurried shoppers appear to be a profitable group to target. For future stores, planner might want to pay extra attention to candidate neighbourhoods employment, income and household size demographics.

References

Code

```
# load necessary packages
library("pander") # for tables
library("glmnet")
```

```

# Load supermarket data from github
githubURL = "https://github.com/jakob-ra/Supervised-Machine-Learning/raw/master/HW_2/supermarket1996.RD
load(url(githubURL))
df = subset(supermarket1996, select = -c(STORE, CITY, ZIP, GROCCOUP_sum, SHPINDX) )
attach(df)

# create summary statistics
pander(summary(supermarket1996) , caption = 'Summary statistics')

# Create vector y (turnover)
y = GROCERY_sum

# Create matrix X of predictor variables
X = subset(df, select = -GROCERY_sum)

# Standardize variables
X = scale(X)
y = scale(y)

# Loss function for MM-algorithm
loss = function(y,X,alpha,lambda,beta,epsilon){
  # Loss function of MM-algorithm given the data, parameters beta, hyperparameters lambda and alpha,
  # and convergence threshold epsilon. Returns both the value of the loss function and the D matrix

  n = length(y) # Number of observations
  p = length(beta) # Number of parameters

  D = matrix(0,p,p) # Initialize p times p matrix of 0s
  beta_abs_list = rep(0,p) # Vector to hold absolute values of beta (will need them later)

  for (j in seq(0,p,1)){
    beta_abs_list[j] = abs(beta[j])
    D[j,j] = 1/max(c(beta_abs_list[j],epsilon))
  }

  c = 1/(2*n)*t(y)%*%y + 1/2*lambda*alpha*sum(beta_abs_list)
  l = 1/2*t(beta)%*%(1/n*t(X)%*%X + lambda*(1-alpha)*diag(p) +
    lambda*alpha*D)%*%beta - 1/n*t(beta)%*%t(X)%*%y + c

  return(list(l,D))
}

# MM-algorithm for minimizing the elastic net loss function
elastic_net_MM = function(y,X,alpha,lambda,beta_0=rep(0,(dim(X)[2])),epsilon=10^(-8)){
  # Fits an elastic net model via MM-algorithm. Returns vector beta_hat for given data,
  # hyperparameters lambda and alpha, initial parameter guess beta_0 (default is a vector of 0s)
  # and convergence threshold epsilon (default is 10^(-8)).

  n = length(y) # Number of observations
  p = (dim(X)[2]) # Number of parameters

  # Set values for the first iteration
  beta = beta_0

```

```

k = 1
# Set l_new and l_old so that (l_old-l_new)/l_old > epsilon is true in the first iteration
l_new = 0
l_old = 1

# Iterate the approximation until convergence
while ((l_old-l_new)/l_old > epsilon){
  l_and_D = loss(y,X,alpha,lambda,beta,epsilon) # Loss returns both the value of the loss function and the matrix D
  l_old = l_and_D[[1]] # Gets value l of loss function
  D = l_and_D[[2]] # Gets matrix D, which we already computed in the loss function
  A = 1/n*t(X)%*%X + lambda*(1-alpha)*diag(p) + lambda*alpha*D
  beta = solve(A, 1/n*t(X)%*%y)
  l_new = loss(y,X,alpha,lambda,beta,epsilon)[[1]]

  # print(k)
  k = k + 1
  #print(l_old-l_new)
}

return(beta)
}

k_fold_crossval = function(y,X,k,alpha,lambda){
  # Function for k-fold crossvalidation, returns RMSE for given alpha and lambda
  n = length(y) # Number of observations
  p = (dim(X)[2]) # Number of parameters

  # Reshuffle data and split into k groups of equal size
  reshuffled_indices = sample(seq(1,n,1), n, replace=FALSE) # Shuffle indices of our n observations
  n_test = floor(n/k) # n=77 and k=10 would give 7 obs. in the test set the rest of the observations
  # in the training set. For n=77, using k=11 or k=7 is advisable to get evenly sized groups.

  MSE=array(0,k) # Vector to hold the MSEs for each fold

  # Loop over the k folds and save MSEs
  for (i in seq(1,k,1)){
    # Divide data into test and training
    y_test = y[c(reshuffled_indices[(1+(i-1)*n_test):(i*n_test)])]
    y_train = y[-c(reshuffled_indices[(1+(i-1)*n_test):(i*n_test)])]
    X_test = X[c(reshuffled_indices[(1+(i-1)*n_test):(i*n_test)])]
    X_train = X[-c(reshuffled_indices[(1+(i-1)*n_test):(i*n_test)])]

    beta = elastic_net_MM(y_train,X_train,alpha,lambda) # get beta estimates for training data from MM
    fitted_val = (X_test)%*%beta # get y_hat for test data
    MSE[i]=1/n_test*t(y_test-fitted_val)%*%(y_test-fitted_val) # save MSE values for test data
  }

  RMSE = sqrt(sum(MSE)/k)

  return(RMSE)
}

```

```

min_RMSE = function(y,X,k=7,alpha_values=seq(0,1,length.out=3),lambda_values=10^seq(-3, 4, length.out =
  # Tunes hyperparameters alpha and lambda via k-fold crossvalidation.
  # Returns optimal combination of alpha and lambda
  # as well as the resulting RMSE.

  RMSE = matrix(0,length(alpha_values),length(lambda_values)) # create matrix to hold RMSE for each
                                                                # hyperparameter combination

  # Loop over hyperparameter combinations
  for (i in 1:length(alpha_values)){
    for (j in 1:length(lambda_values)){
      RMSE[i,j] = k_fold_crossval(y,X,k,alpha_values[i],lambda_values[j]) # Fill the matrix with RMSE
    }
  }

  min_index = arrayInd(which.min(RMSE), dim(RMSE)) # Find index of lowest RMSE
  alpha_min = alpha_values[min_index[1]] # alpha value at lowest RMSE
  lambda_min = lambda_values[min_index[2]] # lambda value at lowest RMSE
  RMSE_min = RMSE[min_index] # lowest RMSE

  return(list(alpha_min,lambda_min,RMSE_min))
}

## We want to compare our result with glmnet function, which cannot cross-validate alpha
## -> compare for fixed alpha
alpha = 0.5

# Find best cross validated lambda from glmnet, letting the function choose
# the sequence of lambdas to try
result.cv <- cv.glmnet(X, y, alpha=alpha, folds=7)
print(result.cv$lambda.min)
# Note that the result here varies quite a bit but is frequently equal or close to 0.05

# Find best cross validated lambda from our own function,
# with a fine sequence of lambdas between 0 and 1
print(min_RMSE(y,X,alpha_values=alpha,lambda_values=seq(0.01,1,0.01))[[2]])
# --> Results are equal or very close

## Knowing that everything works as intended, find the best combination of alpha and lambda
result = min_RMSE(y,X,alpha_values=seq(0,1,0.1),lambda_values=seq(0.01,1,0.01)) # This takes a while
print(result)
# Get point estimates on full sample for the optimal values of alpha and lambda
beta_hat = elastic_net_MM(y,X,0.6,0.04)

# set very low point estimates (rounding errors) to 0
zero_ind = arrayInd(which(beta_hat < 10^(-5)), dim(beta_hat))
beta_hat[zero_ind] = 0
pander(t(beta_hat[beta_hat[,1]!=0,]) ,
  caption = "Non-zero point estimates for (lambda,alpha) = (0.04,0.6)")

# This prints the code chunks at the end

```